

Introduction to Deep Learning

Tutorial 2

Gabriel Deza

Department of Industrial Engineering
Tel Aviv University

November 11th & 13th, 2025

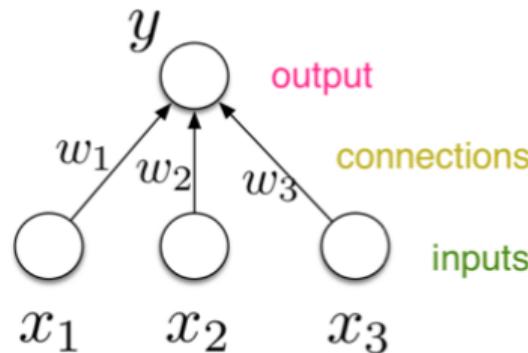
Overview

Last tutorial: review of linear models, gradient descent, PyTorch.

Today:

- Multilayer Perceptron
- Activation functions
- Backprop
- Regularization techniques

Limits of linear models



An equation for a linear model with bias:

$$y = \phi(\mathbf{w}^\top \mathbf{x} + b)$$

Annotations explain the components:

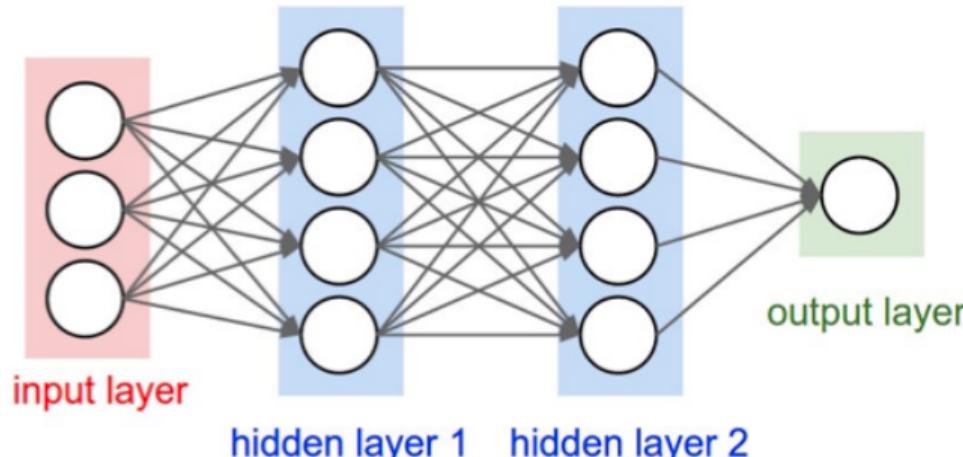
- "output" points to the y variable.
- "weights" points to the term $\mathbf{w}^\top \mathbf{x}$.
- "bias" points to the term b .
- "activation function" points to the ϕ symbol.
- "inputs" points to the \mathbf{x} variable.

Can view logistic and linear regression as a single unit.

In lecture, we saw it was not possible to model **XOR** as it is not linearly separable.

Connecting many such units in a network → more expressive power.

Multilayer Perceptron

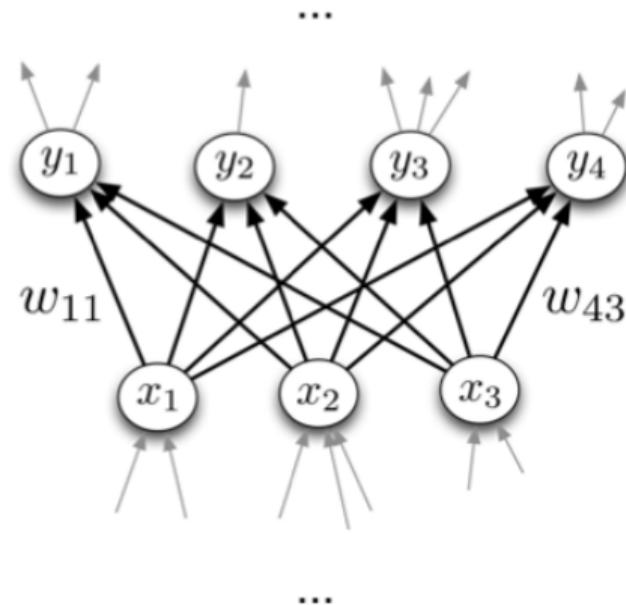


- Directed acyclic graph (DAG)
- Also known as **feed forward** neural network (no cycles)
- Layers:
 - 1 **input** layer to feed in input features.
 - h **hidden** layers to compute advanced features.
 - 1 **output** layer to produce interpretable output.

- Each layer connects N input units to M output units. When all input units are connected to output units → **Fully connected layer**.
- Inputs/outputs of a layer is not the same as inputs/outputs to the network.

$$\mathbf{y} = f(\mathbf{x}) = \phi(\mathbf{Wx} + \mathbf{b})$$

- Why do we need activation functions ϕ ? What if $\phi(x) = x$?
- Multilayer feed-forward neural nets with nonlinear activation functions are **universal approximators**: they can approximate any function arbitrarily well.



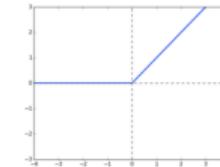
Activation Functions

- Non-linear activation functions allow us to go beyond an affine transform.
- Multilayer feed-forward neural nets with nonlinear activation functions are **universal approximators**: they can approximate any continuous function arbitrarily well.



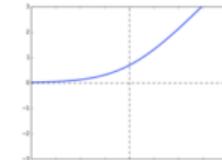
Linear

$$y = z$$



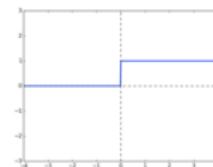
Rectified Linear Unit
(ReLU)

$$y = \max(0, z)$$



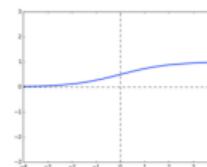
Soft ReLU

$$y = \log(1 + e^z)$$



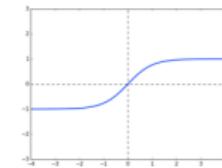
Hard Threshold

$$y = \begin{cases} 1 & \text{if } z > 0 \\ 0 & \text{if } z \leq 0 \end{cases}$$



Logistic

$$y = \frac{1}{1 + e^{-z}}$$



Hyperbolic Tangent
(tanh)

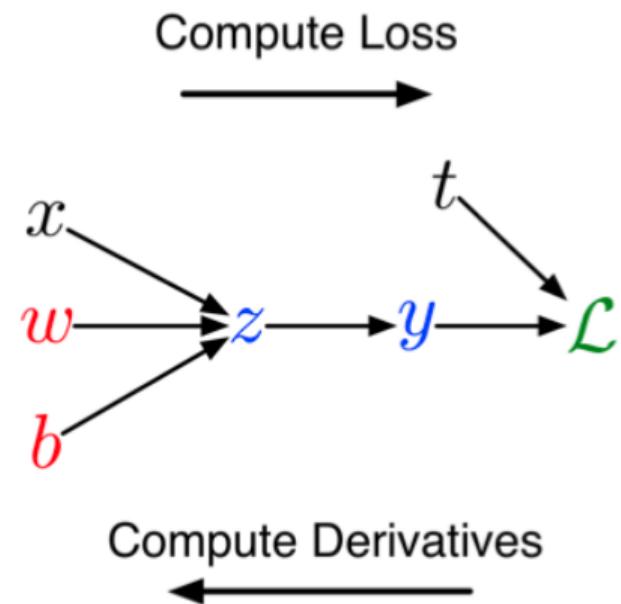
$$y = \frac{e^z - e^{-z}}{e^z + e^{-z}}$$

Backpropagation

- Gradient descent (GD) updates parameters using the cost gradient:

$$\mathbf{W} \leftarrow \mathbf{W} - \alpha \nabla \mathcal{J}(\mathbf{W})$$

- The cost gradient is the average of $\frac{d\mathcal{L}}{d\mathbf{w}}$ across all training examples.
- Backpropagation** efficiently computes these gradients in neural networks. It relies on:
 - Chain rule
 - Computation graphs



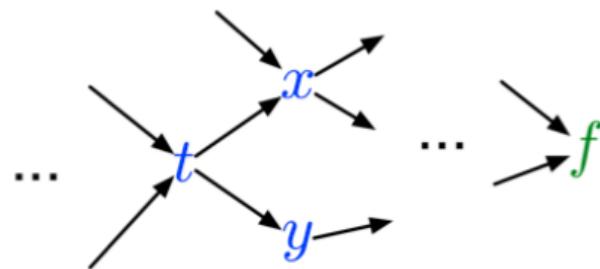
Backpropagation

- Suppose we have a function $f(x, y)$ and functions $x(t)$ and $y(t)$. (All the variables here are scalar-valued.) Then

Mathematical expressions
to be evaluated

$$\frac{df}{dt} = \frac{\partial f}{\partial x} \frac{dx}{dt} + \frac{\partial f}{\partial y} \frac{dy}{dt}$$

Values already computed
by our program



Backpropagation

Full backpropagation algorithm:

Let v_1, \dots, v_N be a **topological ordering** of the computation graph (i.e. parents come before children).

v_N denotes the variable we're trying to compute derivatives of (e.g. the loss).

forward pass

For $i = 1, \dots, N$: Compute v_i as a function of $\text{Pa}(v_i)$

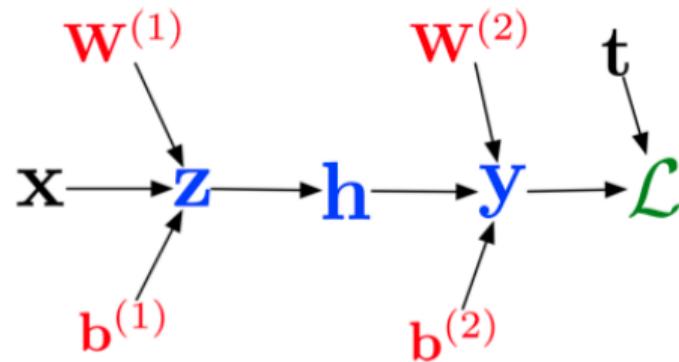
backward pass

$$\bar{v}_N = 1$$

For $i = N - 1, \dots, 1$: $\bar{v}_i = \sum_{j \in \text{Ch}(v_i)} \bar{v}_j \frac{\partial v_j}{\partial v_i}$

MLP Example

MLP example in vectorized form:



Forward pass:

$$z = \mathbf{W}^{(1)}\mathbf{x} + \mathbf{b}^{(1)}$$

$$\mathbf{h} = \sigma(z)$$

$$\mathbf{y} = \mathbf{W}^{(2)}\mathbf{h} + \mathbf{b}^{(2)}$$

$$\mathcal{L} = \frac{1}{2}\|\mathbf{t} - \mathbf{y}\|^2$$

Backward pass:

$$\bar{\mathcal{L}} = 1$$

$$\bar{\mathbf{y}} = \bar{\mathcal{L}}(\mathbf{y} - \mathbf{t})$$

$$\overline{\mathbf{W}^{(2)}} = \bar{\mathbf{y}}\mathbf{h}^\top$$

$$\overline{\mathbf{b}^{(2)}} = \bar{\mathbf{y}}$$

$$\bar{\mathbf{h}} = (\mathbf{W}^{(2)})^\top \bar{\mathbf{y}}$$

$$\bar{\mathbf{z}} = \bar{\mathbf{h}} \odot \sigma'(z)$$

$$\overline{\mathbf{W}^{(1)}} = \bar{\mathbf{z}}\mathbf{x}^\top$$

$$\overline{\mathbf{b}^{(1)}} = \bar{\mathbf{z}}$$

Putting it all together

- Initialize model parameters \mathbf{W} randomly*

- **Batch Gradient Descent (GD):**

- At iteration k , compute the cost (average loss) over the **entire dataset**.

$$\mathbf{W}^{(k+1)} = \mathbf{W}^{(k)} - \alpha \nabla \mathcal{J}(\mathbf{W}^{(k)})$$

- **Stochastic Gradient Descent (SGD):**

- Randomly select a **mini-batch** (small subset of the training data).
 - At iteration k , perform GD update with cost averaged over the **mini-batch**. Update.
 - **Intuition:** noisy approximation of the full gradient.
 - **Tip:** subsample without replacement.
 - **Epoch:** enough iterations to cover the entire dataset once.

* In practice, use smarter initialization (e.g. Xavier, He).

Code: Training MLP using PyTorch

Notebook time :)

https://colab.research.google.com/drive/1DaOKVM_t5S4zXCHY--_OHFAdqbye2IwN?usp=sharing