

Introduction to Deep Learning

Tutorial 4

Gabriel Deza

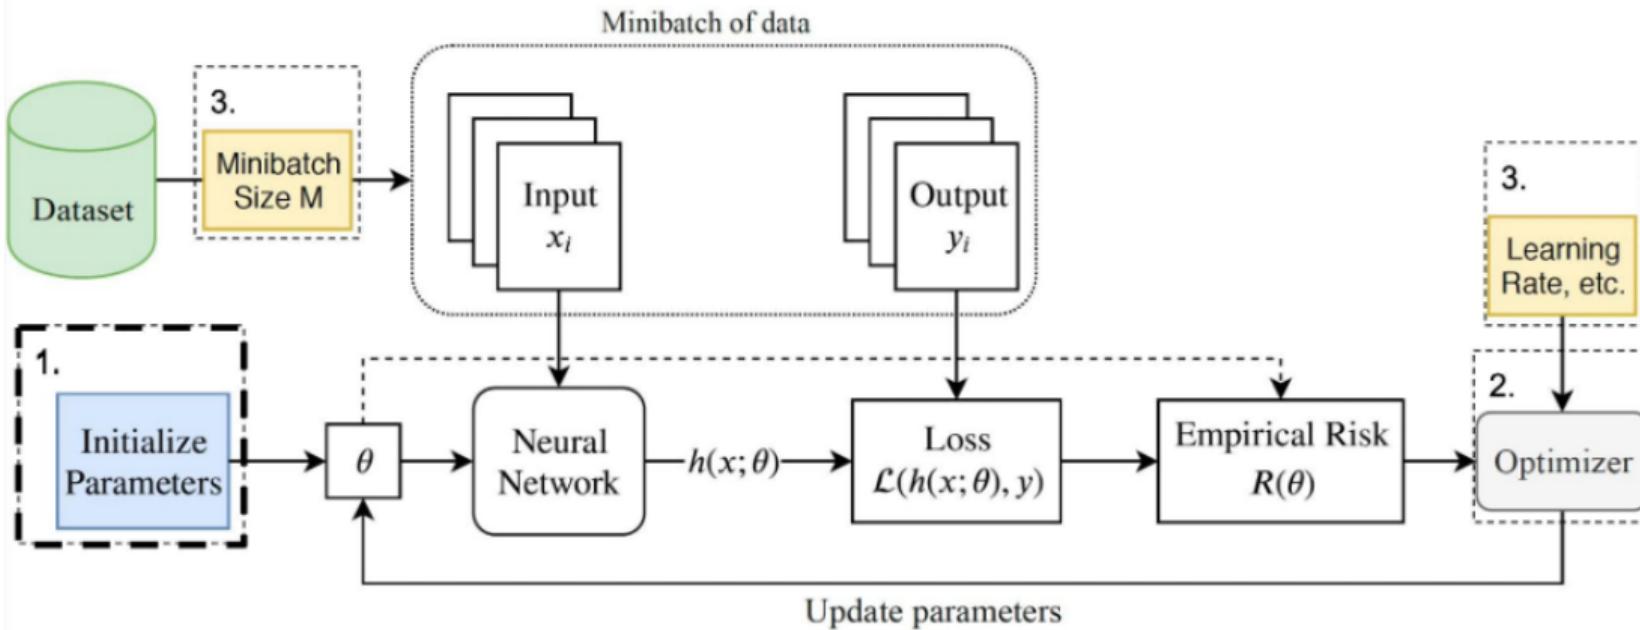
Department of Industrial Engineering
Tel Aviv University

November 24, 2025

Overview

- Overall training loop
 - Initialization
 - Optimization:
 - Gradient Descent
 - Momentum, Nesterov accelerated Momentum
 - Learning rate schedulers: Adagrad, RMSProp, Adam
 - Hyperparameter tuning

Neural Network Training Loop



Initializing weights

TLDR: The art of weight initialization is already largely solved.

Let PyTorch handle it unless you know what you are doing.

```
1 ✓ import torch # type: ignore
2   import torch.nn as nn # type: ignore
3   import torch.nn.functional as F # type: ignore
4
5 ✓ class MLP(nn.Module): #every NN module should inherit from nn.Module
6   ✓ def __init__(self, input_dim, hidden_dim, output_dim):
7     super(MLP, self).__init__() # initialize parent nn.Module
8     # Define layers
9     self.fc1 = nn.Linear(input_dim, hidden_dim) # Input → Hidden
10    self.fc2 = nn.Linear(hidden_dim, hidden_dim) # Hidden → Hidden
11    self.fc3 = nn.Linear(hidden_dim, output_dim) # Hidden → Output
12
13  ✓ def forward(self, x):
14    h1 = F.relu(self.fc1(x))
15    h2 = F.relu(self.fc2(h1))
16    y = self.fc3(h2) # final layer usually left without activation (depends on task)
17    return y
```

Optimization at 3000 feet

- **TLDR:** use Adam.
- If you have time, try different optimizers/hyperparameters.

```
# Model, loss, optimizer
model = MLP(input_dim, hidden_dim, output_dim).to(device)
criterion = nn.MSELoss()
#GD or SGD
optimizer = torch.optim.SGD(model.parameters(), lr=lr)
# SGD with momentum
optimizer = torch.optim.SGD(model.parameters(), lr=lr, momentum=0.9)
# RMSprop
optimizer = torch.optim.RMSprop(model.parameters(), lr=lr)
# Adagrad
optimizer = torch.optim.Adagrad(model.parameters(), lr=lr)
# Adam
optimizer = torch.optim.Adam(model.parameters(), lr=lr)
```

Optimization formally

- Given a training set: $\{(x_1, y_1), \dots, (x_n, y_n)\}$
- Prediction function: $h(x; \theta)$
- Define a loss function: $\mathcal{L}(h(x; \theta), y)$
- Find the parameters: $\theta = (\theta_1, \dots, \theta_k)$ which minimizes the **empirical risk** $R(\theta)$:

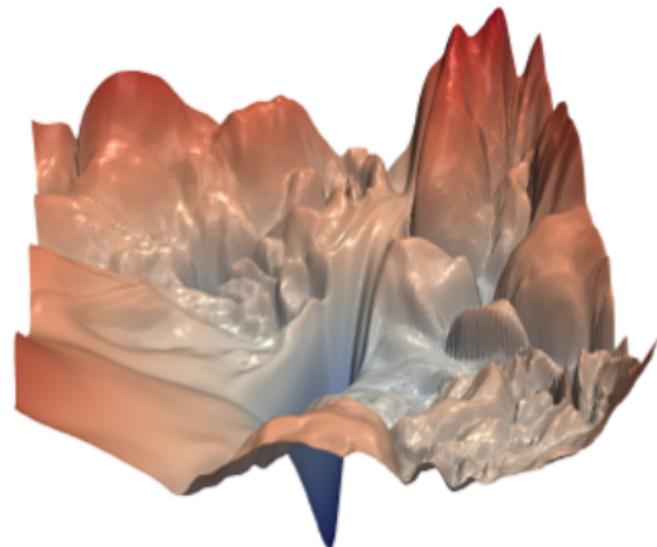
$$\min_{\theta} R(\theta) = \min_{\theta} \frac{1}{n} \sum_{i=1}^n \mathcal{L}(h(x_i; \theta), y_i)$$

Optimum satisfies $\nabla R(\theta^*) = 0$

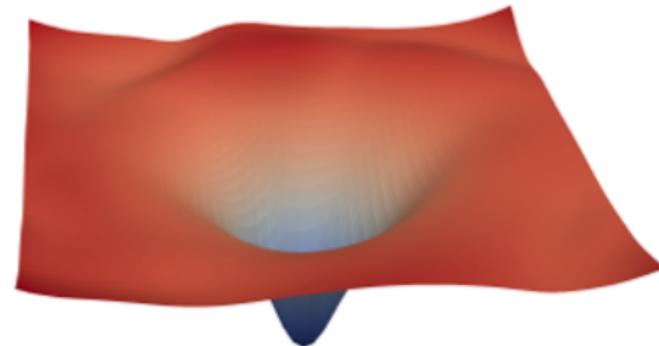
Skip Connections (Residuals)

Skip Connections: $x_{l+1} = F_l(x_l) \rightarrow x_{l+1} = x_l + F_l(x_l)$

- Learn a residual correction instead of a full mapping.
- Identity path helps gradients flow \Rightarrow easier deep training.



(a) without skip connections



(b) with skip connections

Gradient Descent

Geometric view:

- Gradient is orthogonal to the level set.
- The negative gradient gives the steepest descent direction.

Update rule:

$$\theta^{(k+1)} = \theta^{(k)} - \eta \nabla R(\theta^{(k)})$$

SGD: Uses random mini-batches to approximate the gradient. (GD is SGD with the full dataset as one batch.)

```
from torch.utils.data import DataLoader  
  
train_dataloader = DataLoader(training_data, batch_size=64, shuffle=True)  
test_dataloader = DataLoader(test_data, batch_size=64, shuffle=True)
```

Gradient Descent with Momentum

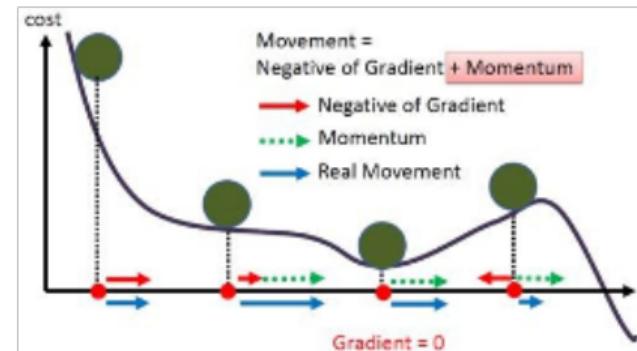
- Initialize parameters $\theta^{(0)}$ and momentum $\delta^{(0)} = \mathbf{0}$.
- For $k = 0, 1, 2, \dots$ until convergence:

- **Update momentum:**

$$\delta^{(k+1)} = -\eta \nabla R(\theta^{(k)}) + \alpha \delta^{(k)}$$

- **Update parameters:**

$$\theta^{(k+1)} = \theta^{(k)} + \delta^{(k+1)}$$



- **Pros:** accelerates learning by accumulating a velocity from past gradients; helps damp oscillations. Use $\alpha = 0.9$.

AdaGrad Optimizer

- Initialize parameters $\theta^{(0)}$ randomly and initialize accumulator $G^0 = 0$
- For each iteration $k = 0, 1, 2, \dots$:
 - Accumulate squared gradients:

$$G_i^{(k+1)} = G_i^{(k)} + \nabla R(\theta_i^{(k)})^2$$

- Update parameters:

$$\theta_i^{(k+1)} = \theta_i^{(k)} - \frac{\eta}{\sqrt{G_i^{(k)}} + \epsilon} \odot \nabla R(\theta_i^{(k)})$$

- G is a vector of how much each parameter has changed over time.
- **Intuition:** increases the learning rate for sparse features and decreases it for frequent features, based on gradient history.

RMSProp/Adadelta Optimizer

- Maintain an exponential moving average of squared gradients:

$$G_i^{(k+1)} = \gamma G_i^{(k)} + (1 - \gamma) \nabla R(\theta_i^{(k)})^2$$

- Parameter update:

$$\theta_i^{(k+1)} = \theta_i^{(k)} - \frac{\eta}{\sqrt{G_i^{(k+1)}} + \epsilon} \nabla R(\theta_i^{(k)})$$

- Intuition:** Like AdaGrad, adapts learning rates per parameter, but emphasizes recent gradients ($\gamma \approx 0.9$).
- Pros:** Prevents learning rates from vanishing; stabilizes training.

Adam Optimizer

- Maintain exponential moving averages of gradients and squared gradients:

$$m_k = \beta_1 m_{k-1} + (1 - \beta_1) \nabla R(\theta^{(k)}), \quad v_k = \beta_2 v_{k-1} + (1 - \beta_2) (\nabla R(\theta^{(k)}))^2$$

- Apply bias correction:

$$\hat{m}_k = \frac{m_k}{1 - \beta_1^k}, \quad \hat{v}_k = \frac{v_k}{1 - \beta_2^k}$$

- Parameter update:

$$\theta_{k+1} = \theta_k - \frac{\eta}{\sqrt{\hat{v}_k} + \epsilon} \hat{m}_k$$

- Intuition:** Combines Momentum (via m_k) and RMSProp (via v_k). Adaptive learning rates + bias correction = fast and stable training.
- Defaults:** $\beta_1 = 0.9$, $\beta_2 = 0.999$, $\epsilon = 10^{-8}$.

[Adam: A method for stochastic optimization](#)

D Kingma, J Ba

International Conference on Learning Representations

235387

2015

[Layer normalization](#)

J Ba, JR Kiros, GE Hinton

Advances in NIPS 2016 Deep Learning Symposium, arXiv preprint arXiv:1607.06450

17071

2016

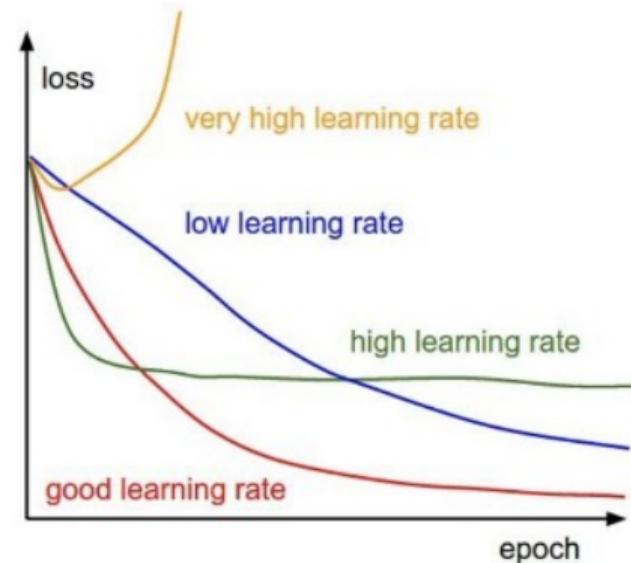
Learning Rate

Ideal Learning Rate should be:

- Should not be too big (objective will blow up)
- Should not be too small (takes longer to converge)

Convergence criteria:

- Change in objective function is close to zero
- Gradient norm is close to zero
- Validation error starts to increase (early-stopping)



Idealized cartoon depiction of different learning rates.

Image credit: Andrej Karpathy

Learning Rate Scheduling: Brief Overview

Anneal (decay) learning rate over time so the parameters can settle into a local minimum.

Typical decay strategies:

1. **Decay:** reduce the learning rate somehow monotonically over time (many possible ways).
2. **Cosine Scheduling:** decay slowly, then quickly, then slowly again. Not clear why this works well.
3. **Schedule-Free Tuning:** uses clever math to tune the learning rate automatically.

Batch Size

Definition: Number of training examples used to compute the gradient at each iteration.

- Typical small batch sizes are powers of 2: 32, 64, 128, 256, 512
- Large batches are in the thousands

TLDR: Pick largest batch size your GPU/TPU setup supports

Large Batch Size effects:

- Fewer parameter updates per epoch
- More accurate gradient estimate
- Better parallelization efficiency \Rightarrow faster wallclock training
- **May hurt generalization** (risk of sharper minima or poorer local optima)

Architecture Choice

- **Top-level rule:** Start with a simple, proven architecture. Add complexity only if needed.
- Avoid chasing every “new thing” unless it clearly helps.
- With experience, you’ll develop intuition about which techniques work best for particular tasks.

Common Bells and Whistles:

- Layer Normalization
- Batch Normalization
- Dropout
- Residual Connections
- Extensive hyperparameter tuning

```
class MLP(nn.Module): #every NN module should inherit from nn.Module
    def __init__(self, input_dim, hidden_dim, output_dim):
        super(MLP, self).__init__() # initialize parent nn.Module
        # Define layers
        self.fc1 = nn.Linear(input_dim, hidden_dim) # Input → Hidden
        self.fc2 = nn.Linear(hidden_dim, hidden_dim) # Hidden → Hidden
        self.fc3 = nn.Linear(hidden_dim, output_dim) # Hidden → Output
        self.dropout = nn.Dropout(p=0.5) # Dropout regularization
        self.bn1 = nn.BatchNorm1d(hidden_dim) # Batch normalization after fc1
        self.bn2 = nn.BatchNorm1d(hidden_dim) # Batch normalization after fc2

    def forward(self, x):
        h1 = F.relu(self.bn1(self.fc1(x)))
        h1 = self.dropout(h1)
        h2 = F.relu(self.bn2(self.fc2(h1)))
        h2 = self.dropout(h2)
        y = self.fc3(h2) # output layer (no activation here)
        return y
```

Hyperparameter Optimization (HPO)

- Many settings (LR, batch size, weight decay, depth, dropout, scheduler) strongly affect performance.
- HPO frameworks (e.g., **Optuna**, Ray Tune, Hyperopt, SMAC) automate this search.
- They run multiple *trials*: train the model with different hyperparameters and measure validation score.
- A smart search strategy balances *exploration* (try new regions) and *exploitation* (refine good ones).
- Often includes *early stopping/pruning* to kill bad trials fast and save compute.
- Visualizing optimizers (Ruder blog)